


Runtime (all 60 cases)



Agent median 206.7s vs Exhaustive 19.4s (ratio 10.6×). Agent wall-clock includes local LLM inference (qwen3:4b) and is not implied by tool-call count.